
Unsupervised Deep Audio Denoising Algorithms via Hybrid Reinforcement Learning and Genetic Algorithm Meta-Learning

A PREPRINT

Julian M. Kleber 

Email: julian.m.kleber@gmail.com

19 July 2026

Code (ReelSynth engine + DenoiseOpt meta-search): <https://github.com/reeldemo/reelsynth>,
<https://github.com/reeldemo/denoise-opt-meta>

ABSTRACT

When a single-period wavetable (or any short cyclic buffer) is played with phase wrap, a sample discontinuity at the period boundary injects a broadband transient that repeats at the fundamental. We treat this cycle-local seam crackle as an audio denoising problem: seam-local operators (DualCosine crossfade, polish, FIR blends, optional tiny residual cells) are baked into the period, then judged by a prolonged residual $R \in [0, 1]$ against an ideal multi-period tiling that withholds the open-wrap cliff (1 = best). We call the residual-scored hybrid search **DenoiseOpt**. An outer hybrid meta-search interleaves GA tournament updates with proximal policy optimization (PPO) architecture proposals, optional PBT exploit–mutate schedules, and MoE soft gates over heterogeneous bake cells. Candidates need not be paired with studio-clean waveforms for every trial. On the frozen canonical sine+cliff holdout, a compact favorite reaches $R \approx 0.991$ versus DualCosine $R \approx 0.825$ ($\Delta R \approx +0.166$). A live GPU hybrid campaign has passed 5,000 clean iterations with champion $R \approx 0.9909$ against DualCosine ≈ 0.8166 on the runner geometry.

Keywords unsupervised learning · audio denoising · DenoiseOpt · neural architecture search · reinforcement learning · genetic algorithms · wavetable synthesis · wrap discontinuity · meta-learning

1 Introduction

Wavetable oscillators store one period of a waveform and advance a read pointer modulo the table length. If $x[0] \neq x[L-1]$, the wrap is a hard discontinuity. Under cyclic playback it becomes a periodic click whose spectrum is broadband and locked to the fundamental. The same geometry appears whenever a short loop buffer is tiled for listening or for objective evaluation. We call the repair task *cycle-local seam restoration*: edit a neighborhood of the wrap (and optionally a lightweight residual cell over the period) so that multi-period playback matches an ideal reference that never opens the wrap cliff. That outer agreement is what we score as audio denoising for this periodic artifact class.

Figure 1 shows one real sine-wrap edge case from the frozen canonical holdout (same `make_batch` tile, seed 20,260,719). Panel (a) is the ideal multi-period sine. Panel (b) is the cracked engine tile after open-wrap cliff modulation ($R \approx 0.950$), with the ideal overlaid so the seam jump is visible. Panel (c) applies DualCosine on that same tile ($R \approx 0.603$ on this severe cliff). Panel (d) is the top DenoiseOpt favorite (champion_iter_000235, batch $R \approx 0.991$) on the same tile ($R \approx 0.9917$). Panels (e)–(f) plot the modulation resid-

ual and a seam zoom so the favorite’s correction is readable against DualCosine.

Classical virtual-analog practice already supplies seam-local tools. Alias-free BLIT/BLEP-style synthesis [1], Poly-BLEP / fractional-delay antialiasing [3], and BLAMP corner corrections [2] reduce wrap and geometric discontinuity energy in-band. Those operators act locally. They leave open which bake parameters and which residual cell, if any, minimize prolonged tiled error against a no-cliff ideal. Label-free restoration work such as Noise2Noise shows that useful reconstructors can be ranked without paired clean targets when the evaluation geometry is explicit [4]. Speech-domain Noise2Noise denoisers extend that idea to audio waveforms [5]. Mixture-invariant unsupervised separation likewise avoids paired clean stems for ranking [29]. On the search side, NAS surveys [16], RL controllers [17, 18], aging evolution [20, 21], population-based hyperparameter schedules [22], CMA-ES tutorials [24], and evolution-guided policy gradients [23] motivate a hybrid outer loop: GA for population diversity, PPO for discrete architecture edits [19], and PBT-style exploit–mutate on elites. Platform-aware and progressive NAS add compute-aware search priors [27, 28]. MoE soft gates suggest routing among heterogeneous tiny cells under one residual score [25]. Waveform separation nets (Wave-




Figure 1: Sine-wrap modulation edge case from the canonical eval holdout (same tile). Ideal reference, cliff-modulated engine (ideal overlay), DualCosine bake, and top meta favorite ($R \approx 0.9917$ on this tile, batch mean $R \approx 0.991$). Bottom: modulation residual and seam zoom. Prolonged residual R is printed per panel.

U-Net, Conv-TasNet, dual-path RNN, SepFormer) and audio DL surveys inform the searchable cell family at bake-scale width [7–9, 11, 12, 30, 31], with U-Net / ResNet / attention primitives as block ancestors [10, 13–15].

The contribution is a meta-learning protocol for this periodic instance. An outer RL+GA(+PBT) loop proposes denoising operator graphs and continuous bake parameters. Each candidate is scored by a prolonged residual $R \in [0, 1]$ ($1 = \text{best}$) between engine-baked multi-period audio and a procedural ideal tiling. Paired studio-clean wavetables are not required for every trial. Differentiable DSP modules [6] place DualCosine / FIR / MLP seam ops in an interpretable stack. Bayesian HPO [26] informs prior families that compete under the same R . A CUDA overnight campaign (PPO+GA+PBT+NAS, depth and MoE biases) has passed a 5,000-iteration clean gate with champion $R \approx 0.9909$ versus DualCosine ≈ 0.8166 on the runner geometry. The frozen canonical holdout and inference benches in Results are the measured evaluation matrices for this report.

Contributions.

1. **Problem framing.** Cast wavetable wrap discontinuity / cycle-local seam restoration as audio denoising under prolonged tiled evaluation, separating seam-local diagnostics from residual R .
2. **Hybrid DenoiseOpt outer loop.** Specify RL+GA(+PBT) search over discrete operator graphs and continuous bake parameters, with named branches (GA crossover, PPO mutation, PBT exploit, MoE soft gates).
3. **Literature protocol.** Separate used versus screened prior work for audio denoise cells and RL–evolutionary hybrids, grounded on downloaded source PDFs.
4. **Open evaluation path.** Release companion code for the synthesizer engine and meta search, with frozen 5k-gate figures and matched classical-vs-AI benches.

Section 2 organizes related work by mechanism. Sections 3–4 outline the residual objective and search protocol. Results report the 5k-gate freeze and inference benches.

2 Related Work

We organize prior work by mechanism. **Used** lines are methodological or comparative anchors for residual-scored RL+GA meta-search. **Screened** lines were inspected as contrast but are not treated as dependencies. Paraphrases prefer attached abstracts and PDF snippets over title-only citation.

2.1 Discontinuity control and modular audio DSP

Alias-free digital synthesis of classic analog waveforms (BLIT/BLEP lineage) states the wrap discontinuity problem for trivial oscillators [1]. PolyBLEP-style fractional-delay corrections give practical antialiasing oscillators for subtractive synthesis [3]. BLAMP corner rounding treats discontinuities in the first derivative (triangular edges, clipping, rectification) without heavy oversampling [2]. DDSP reframes DSP modules as composable blocks inside learning pipelines [6]. We use that framing for small bake/FIR/MLP seam cells. We do not train end-to-end neural vocoders for every meta trial. These works are **used** as domain anchors for the operator vocabulary and for why prolonged cyclic playback is the outer judge.

2.2 Label-free restoration

Noise2Noise shows that restoration mappings can be learned from corrupted observations alone by exploiting structure in the corruption process [4]. Speech Noise2Noise denoisers apply the same idea to waveforms without clean paired audio [5]. Mixture invariant training ranks unsupervised separation systems without stem labels [29]. That philosophy motivates ranking denoise candidates without paired studio-clean wavetable cycles. The residual compares an engine-baked tiling to a procedural ideal that withholds open-wrap cliffs. We **use** this family as conceptual justification for un-

supervised residual ranking. We do not retrain a Noise2Noise CNN on images.

2.3 Waveform denoising and separation cells

Deep learning for audio surveys the move from hand-designed features to learned time/frequency models [7]. Wave-U-Net performs end-to-end source separation with multi-scale 1-D U-Nets [8]. Improved Wave-U-Net speech enhancement and singing-voice U-Nets popularized encoder-decoder skips on audio [9, 10]. The original U-Net, ResNet residuals, and Transformer attention supply the block primitives we shrink to bake width [13–15]. Conv-TasNet and dual-path RNN demonstrated strong time-domain separation with temporal convolutions and long-context paths [11, 12]. SepFormer and audio spectrogram transformers push attention-heavy audio stacks further [30, 31]. These lines are **used** as design priors for tiny searchable cells (shallow U-Net / dilated / dual-path / attention blocks under a strict per-trial bake budget). They are not frozen full-scale separators inside every meta trial.

2.4 NAS and RL controllers

Elsken et al. taxonomize NAS by search space, strategy, and evaluation cost [16]. Their performance-estimation axis includes lower-fidelity proxies such as shorter training, data subsets, and lower-resolution images [16]. Zoph and Le couple RL controllers to architecture search [17]. ENAS adds parameter sharing for efficient discrete search [18]. Progressive and platform-aware NAS add staged evaluation and latency-aware rewards [27, 28]. We **use** this family as the ancestor of RL proposing architecture edits (add/remove/mutate ops, toggle residual-primary, mutate MLP/FIR), scored with prolonged residual. PPO supplies the clipped-surrogate policy update when the overnight loop logs PPO mutation branches [19].

2.5 Evolutionary NAS and population schedules

Large-scale and regularized / aging evolution showed that evolutionary NAS can discover competitive image classifiers when population turnover is carefully regularized [20, 21]. Population Based Training jointly adapts weights and hyperparameters in an asynchronous population [22]. Practical Bayesian optimization [26] and CMA-ES tutorials [24] inform local Bayesian and continuous evolutionary prior families. These are **used** for GA tournament/crossover branches, PBT-style exploit-mutate, and literature-combo priors, implemented at synthesizer-trial scale.

2.6 Evolutionary reinforcement learning hybrids

Khadka and Tumer’s evolution-guided policy gradient (ERL) interleaves an evolutionary population with an RL actor to mitigate sparse reward and brittle hyperparameters in deep RL [23]. We **use** ERL as the spirit of interleaving GA population steps with PPO updates at tiny population sizes. We do not claim a full ERL robotics stack.

2.7 Mixture-of-experts soft gating

Sparsely gated MoE layers increase capacity via conditional computation [25]. We **use** MoE as motivation for soft gates

over heterogeneous seam blocks (e.g., U-Net vs dilated vs attention cells) under residual scoring, at bake-scale widths.

2.8 Used versus screened

Used (methodological or comparative anchors).

- Discontinuity / VA lineage: [1–3].
- Unsupervised restoration philosophy: [4, 5, 29].
- Modular / differentiable audio DSP context: [6].
- Audio DL and compact waveform cells: [7–15, 30, 31].
- NAS / RL controllers / PPO: [16–19, 27, 28].
- Evolutionary NAS and population HPO: [20–22, 24, 26].
- ERL hybrid outer loop: [23].
- MoE soft gating prior: [25].

Screened (relevant contrast, not dependencies).

- **MAML** [32]: bi-level / fast-adaptation framing. We do not differentiate through a deep task network. Inner fits are coordinate-style updates on bake parameters. Outer rank is residual R .
- **DARTS** [33]: continuous relaxation of architecture search. We keep discrete ops + small cells under residual scoring.
- **Demucs / Hybrid Demucs / realtime waveform enhancement** [34–36]: strong separators and enhancers, but full multi-scale / hybrid STFT stacks are too heavy per meta trial under overnight bake budgets.
- **DiffWave** [37]: diffusion vocoder. Screened as a generative sampling stack, not a seam bake operator.
- **SEGAN-style adversarial enhancement** [38]: optional tiny adversarial cues only. Full GAN training loops screened for per-trial cost.

Positioning. Relative to this map, the present work contributes (i) a prolonged residual objective for periodic seam denoising, (ii) an explicitly hybrid RL+GA(+PBT) meta-outer loop with honest branch naming, and (iii) a used/screened literature protocol with downloaded source PDFs. We do not claim SOTA on general speech enhancement or music separation corpora. The primary claim is protocol-level: residual-scored meta-search for cycle-local audio denoise operators under a declared compute budget.

3 Methods

Notation. Throughout, $x \in \mathbb{R}^L$ is the stored single-period wavetable (raw engine input). $y = \Theta(x; \theta)$ is the engine-baked cycle after seam ops and optional residual cells with parameters θ . r^* is a procedural multi-period ideal from the same seed with open-wrap cliffs withheld (not a studio-clean recording). $R \in [0, 1]$ is the prolonged residual score defined below (1 = best). These symbols are fixed for the residual objective. They do not follow the Noise2Noise clean/noisy pair convention.

Residual score. With tiling length N (typically $N=16$),

$$R = \text{clamp}\left(1 - \frac{\text{rms}(y_{\text{tiled}} - r_{\text{tiled}}^*)}{\max(\text{rms}(r_{\text{tiled}}^*), \varepsilon)}, 0, 1\right). \quad (1)$$

By construction $R \in [0, 1]$ with 1 best. Soft shape gates (when enabled) down-weight mid-cycle damage relative to seam-local error.

Searchable operator. Candidates combine classical seam ops (DualCosine, polish, soft seam, FIR blends) with tiny residual cells (MLP/FIR/U-Net-lite/attention-lite) and optional MoE soft gates. Bake parameters θ live in a bounded box. Architecture strings are discrete graphs over those ops.

Hybrid outer loop. Each iteration may (i) run GA tournament selection with crossover/mutation, (ii) sample PPO actions that edit ops/hyperparameters, (iii) apply PBT-style exploit–mutate on elites, and (iv) evaluate literature-combo priors. Selection is always by residual R (and logged auxiliaries). The live overnight configuration is tagged PPO+GA+PBT+NAS+depth+MoE with population size 12, plateau adaptation every 1000 iterations, and CUDA scoring on an RTX 3090.

4 Experiments

4.1 Dataset

Creation. The paper-facing evaluation corpus is a frozen draw from the same synthetic generator used by the overnight CUDA runner (`make_batch` in `overnight_gpu_rl_arch.py`). Each cycle has length $L=256$. A two-harmonic sine is drawn with frequency $\sim \mathcal{U}(1, 4)$ and phase $\sim \mathcal{U}(0, 2\pi)$. An open-wrap seam cliff of amplitude $\pm \mathcal{U}(0.08, 0.43)$ is applied over `SEAM_W=8` samples at both ends. Seam-boosted Gaussian noise (scale 0.02 at the ends, $0.15\times$ in the mid-cycle) is added to form the engine input. The ideal withholds the cliff and is used only for residual scoring. Prolonged residual R tiles each cycle $N=16$ times before the RMS ratio. The holdout seed is 20,260,719 (distinct from the overnight search seed 1,902,771,841). Overnight search draws fresh i.i.d. batches from this generator and does not train on the frozen holdout tensors. There is no supervised clean/noisy training split of studio audio. Family labels are absent: this corpus is a single sine+cliff family (Rust `sound_bench` families remain a separate hardness probe in Results).

Dataset metrics. A metrics draw of 4,096 cycles under the holdout seed yields mean ideal RMS 0.721 ± 0.016 , mean engine residual RMS 0.021 ± 0.013 , and mean absolute wrap jump $|x_0 - x_{L-1}|$ of 0.913 ± 0.634 . Identity (no bake) on that draw scores mean $R \approx 0.971$. Method tables below use a matched score batch of 64 cycles from the same seed (tensors persisted under `brand/artifacts/canonical_eval_dataset/`), plus five consecutive seeds for `mean±std`. Figure 1 is the documented sine-wrap edge case: a high-`|wrap|` holdout tile where cliff modulation is visible, scored with DualCosine and the top neural favorite under the same residual R (tile $R \approx 0.9917$, favorite batch mean $R \approx 0.991$).

Overnight hybrid campaign (5k clean gate). A CUDA search tagged PPO+GA+PBT+NAS (depth and MoE biases) runs with a large declared iteration budget (order 10^5). At the 5,000-iteration clean gate the live job reports champion residual $R \approx 0.99093$ against DualCosine baseline ≈ 0.81658 on the same runner geometry ($\Delta R \approx +0.174$). Branch-best residuals at the freeze: GA ≈ 0.99016 , PBT ≈ 0.99012 , PPO ≈ 0.98924 , combo ≈ 0.98854 , NAS ≈ 0.98677 . Figures in Section 5 regenerate from `history.jsonl` at this gate.

Inference and classical matched benches. High- R fitted cells ($R \geq 0.98$) are timed on CUDA (batch 64, prolonged residual). Ten non-learnable bake denoisers share the frozen canonical holdout against the neural favorite (Table 2). Family-conditional hardness uses the separate Rust procedural 100,000-cycle bench reported in Results.

5 Results

Family-conditional hardness (completed bench). On the fitted DenoiseOpt 100,000-cycle bench, denoise score D is lowest for nonlinear ($D \approx 0.39$) and combo ($D \approx 0.40$), then `triple_mix / extreme_overlay` ($D \approx 0.51$ – 0.52), while `harmonic_fft` and `am_fm` sit near $D \approx 0.69$ – 0.70 . Lit-combo champion family stress (prolonged residual) similarly ranks `open_wrap_bias` ($R \approx 0.83$) and `extreme_overlay / combo` ($R \approx 0.88$ – 0.89) below `harmonic_fft` ($R \approx 0.95$). Hard sounds recur.

Overnight hybrid (5k-gate freeze). At $\geq 5,000$ clean iterations the champion residual is $R \approx 0.99093$ with DualCosine ≈ 0.81658 on the overnight runner ($\Delta R \approx +0.174$). Monitoring plots below are regenerated from that freeze.

Inference time at matched score. Across diverse high- R fitted cells ($R \geq 0.98$), we timed CUDA forward passes (batch 64, prolonged residual). Within a tight band of the best residual ($R \geq R_{\max} - 5 \cdot 10^{-4}$), the **favorite** is the lowest-latency model: `champion_iter_000235` with $R \approx 0.9910$, ≈ 3.15 ms/batch, ≈ 37 k parameters (versus a higher-capacity max- R cell at $R \approx 0.9914$ but ≈ 7.18 ms/batch / ≈ 104 k params). Selection rule: maximize R , then minimize latency inside the band.

Table 1: Top-5 distinct fitted architectures on the CUDA inference bench (batch 64). Favorite is latency-best inside the near-max- R band.

Tag	R_{saved}	R_{live}	ms/batch	params
iter_000157	0.9914	0.9913	7.18	104 484
iter_000205	0.9911	0.9907	7.29	70 562
iter_000235*	0.9910	0.9915	3.15	37 489
iter_000397	0.9910	0.9907	6.80	97 262
iter_000229	0.9909	0.9902	6.67	100 500

*Favorite (fastest in $R \geq R_{\text{max}} - 5 \cdot 10^{-4}$ band).

Table 2: Method scores on the canonical sine+cliff holdout (batch 64, seed 20,260,719). Mean $\pm$ std over five consecutive seeds. Identity is a no-op ceiling.

Method	R	R (5-seed)	ms/batch
neural favorite	0.9911	0.9912 ± 0.0002	2.56
identity (no-op)	0.9718	0.9726 ± 0.0021	0.31
seam_fir3	0.9610	0.9616 ± 0.0019	0.67
classic quadratic	0.8449	0.8453 ± 0.0117	1.59
DualCosine	0.8249	0.8258 ± 0.0137	1.35
cosine fade	0.8249	0.8258 ± 0.0137	1.59
crossfade	0.8116	0.8130 ± 0.0145	1.50
hann blend	0.8048	0.8060 ± 0.0156	0.54
soft wide fade	0.7432	0.7378 ± 0.0171	2.64
detrend	0.4079	0.4091 ± 0.0386	0.36
ensemble detrend+DC+FIR	0.4024	0.4027 ± 0.0387	1.56

5.1 Top-5 architectures

Table 1 ranks five distinct fitted bake graphs from the CUDA inference bench (inference_bench.json, batch 64), preferring residual R first and latency on ties, and skipping near-duplicate copies of the same topology signature.

Classical vs AI (canonical holdout). Table 2 scores the classical bake set plus DualCosine and the neural favorite on the frozen corpus of Section 4.1 (CUDA, batch 64, seed 20,260,719). DualCosine reaches $R \approx 0.825$ (≈ 1.35 ms/batch), near the overnight monitor DualCosine ≈ 0.817 on live draws from the same generator. The best *active* classical method is seam-local 3-tap FIR (seam_fir3) at $R \approx 0.961$ / ≈ 0.67 ms/batch (0.962 ± 0.002 over five seeds). Identity (no-op) scores $R \approx 0.972$ and is a ceiling only. The neural favorite reaches $R \approx 0.991$ (0.991 ± 0.0002 over five seeds) at ≈ 2.56 ms/batch ($\approx 37k$ params): $\Delta R \approx +0.030$ vs seam_fir3 and $\Delta R \approx +0.166$ vs DualCosine.

6 Discussion

Where residual mass remains. Across the 100,000-cycle procedural bench and lit-combo family stress tests, wrap error is not uniform. The same generator families repeatedly dominate the residual budget: nonlinear, combo, triple_mix, and extreme_overlay show the weakest denoise / residual scores, while harmonic_fft and am_fm




Figure 2: High- R architectures: residual versus CUDA inference latency (ms/batch). Star: favorite (fastest in the near-max- R band).




Figure 3: Inference latency bars for diverse high- R fitted models.




Figure 4: Canonical holdout: classical bake denoisers vs neural favorite (R vs CUDA ms/batch). Red star: AI favorite.

are comparatively easy. open_wrap_bias can look strong on wrap-energy proxies yet still lag on prolonged residual. Cliff reduction and ideal-matched tiling are related objectives with different rankings.

Implication for this paper. Mean R (and mean DualCosine gaps) therefore compress a structured difficulty map. Reporting only the overnight mean on synthetic sine cliffs understates how much of the remaining $\sim 1\%$ sits in a few hard




Figure 5: Canonical holdout ranking on residual (left) and latency (right). AI in red. DualCosine $R \approx 0.825$.




Figure 6: Champion residual R versus search iteration at the 5k-gate freeze, with DualCosine baseline.




Figure 7: Per-branch best residual trajectories (PPO/RL, GA, PBT, NAS, combo) at the 5k-gate freeze.

families. That structure is a first-class empirical claim of the present work.

Hybrid search vs. plateau. Overnight traces at the 5k gate show early high champions then long plateaus near $R \approx 0.99$. Plateau adaptation (deeper cells, denser MoE graphs, higher GA/PBT weight) is the operational response inside one campaign. It does not by itself rebalance family-conditional risk. A natural next paper is family-conditional and cliff-conditional meta-learning: train or select denoisers per family, or condition the outer loop on measured wrap statistics,

and publish per-family residual curves beside the scalar champion.

Hybrid branches. At the 5k freeze, GA and PBT lead branch-best residuals, with PPO close behind and NAS trailing on this seed.

7 Tooling and AI use

This section discloses research and writing tools, including AI systems, with scope and intent.

Literature discovery. Open literature was gathered with the Klaut Research Gateway (local OpenAlex / Crossref / arXiv providers) under queries spanning virtual-analog anti-aliasing, unsupervised restoration, HPO/PBT, RL-GA hybrids for NAS, and audio denoise architectures. Harvest JSONs live in the companion meta repository. Source PDFs were downloaded for every bibliography entry (artifacts/literature_oa/pdfs/). Paywalled-only items were dropped or replaced.

Paper drafting. Section planning, subsection drafts, and revision passes used Klaut `research_paper_*` tools with local Ollama (qwen3.5:9b) when available, under an explicit scientific workflow (plan $\rightarrow$ write $\rightarrow$ revise $\rightarrow$ export) into an arXiv-style twocolumn template. AI drafts were treated as scaffolding. Claims, numbers, and citations were checked against harvest records, PDF analyses, and measured artifacts. Final prose was edited for DSP terminology and to remove generic filler.

Why AI was used. (1) Scale literature triage across DSP and AutoML without pretending every high-cite hit is in-domain. (2) Accelerate IMRaD scaffolding so human effort focuses on metric honesty (residual vs DualCosine, family hardness, gate vs final overnight claims). (3) Keep an auditable tool chain (plans, PDF snippet analyses, citation flags) over opaque chat-only writing.

Where AI was *not* used. Model fitting, residual scoring, DualCosine baselines, lit-combo / overnight GPU search, and inference-time benchmarks are conventional code (Rust / PyTorch) with deterministic logs and JSON artifacts. AI did not invent experimental numbers.

Other tools. Rust/cargo for the procedural bench and meta binaries. PyTorch CUDA for overnight hybrid search. matplotlib for figures. MiKTeX pdflatex for PDF builds. GitHub for versioned papers and artifacts.

8 Limitations

- Residual uses a procedural ideal (no-cliff sibling), not a perceptual MOS or listening panel.
- Bake metrics are not a substitute for formal audio quality tests.
- The overnight GPU loop currently scores random sine+harmonic seam cliffs. Family-conditional hard-




Figure 8: Full-width overnight monitoring panel at the 5k-gate freeze: champion R (top) and branch-best trajectories (bottom).




Figure 9: Per-trial residual scatter by search branch, with champion overlay (5k-gate freeze).




Figure 10: Champion update timeline. Sparse R labels mark first, peak-mid, and latest improvements to stay legible in twocolumn.

ness is measured on the separate Rust procedural bench. Aligning both evaluators is left to follow-up work.

- Per-trial architecture width is capped for overnight throughput. Full Demucs-diffusion-scale stacks were screened out.

- The 5k-gate freeze reports measured champion and branch residuals on the runner and frozen holdout geometries used in this paper.

- Classical VA/BLEP anchors use author or proceedings PDFs (stilson1996, nam2009polyblep, esqueda2016blamp).

9 Outlook: follow-up publication

A self-contained follow-up can treat *which sounds fail* as the primary object of study: stratified residual heatmaps over the ten procedural families, cliff-size conditioning, and family-specialist or mixture-of-experts denoisers selected by the same RL-GA outer loop. That paper would cite the present mean- R hybrid protocol as the baseline and ask whether beating $R > 0.999$ is mainly an architecture problem or a coverage problem over hard families.

10 Conclusion

We framed wavetable wrap discontinuity as an audio denoising instance and specified DenoiseOpt, a hybrid RL+GA meta-learning outer loop scored by prolonged residual R . Related work separates used anchors (Noise2Noise / speech N2N, PBT, aging evolution, ERL, ENAS/PPO, Wave-U-Net/Conv-TasNet, MoE, BLIT/BLAMP/DPW, DDSP) from screened contrasts (MAML, DARTS, Demucs, DiffWave, full GAN loops). At the 5,000-iteration clean gate the overnight champion reaches $R \approx 0.9909$ versus DualCosine ≈ 0.8166 . A compact inference favorite reaches $R \approx 0.991$ versus DualCosine $R \approx 0.820$ on a matched CUDA bench. Companion code lives at <https://github.com/reeldemo/reel-synth> and <https://github.com/reeldemo/denoise-opt-meta>.

References

- [1] T. Stilson and J. O. Smith. Alias-free digital synthesis of classic analog waveforms. In *Proc. ICMC*, 1996. <https://ccrma.stanford.edu/~stilti/papers/blit.pdf>.
- [2] F. Esqueda, V. Välimäki, and S. Bilbao. Rounding corners with BLAMP. In *Proc. DAFX*, 2016. https://www.dafx.de/paper-archive/2016/dafxpapers/18-DAFx-16_paper_33-PN.pdf.
- [3] J. Nam, V. Välimäki, J. S. Abel, and J. O. Smith. Efficient antialiasing oscillator algorithms using low-order fractional delay filters. In *Proc. DAFX*, 2009. <https://ccrma.stanford.edu/~juhan/pubs/jnam-dafx09.pdf>.
- [4] J. Lehtinen et al. Noise2Noise: Learning image restoration without clean data. In *ICML*, 2018. [arXiv:1803.04189](https://arxiv.org/abs/1803.04189).
- [5] M. M. Kashyap, A. Tambwekar, K. Manohara, and S. Natarajan. Speech denoising without clean training data: A Noise2Noise approach. In *INTERSPEECH*, 2021. https://www.isca-archive.org/interspeech2021/kashyap21_interspeech.pdf.
- [6] J. Engel, L. Hantrakul, C. Gu, and A. Roberts. DDSP: Differentiable digital signal processing. In *ICLR*, 2020. [arXiv:2001.04643](https://arxiv.org/abs/2001.04643).
- [7] H. Purwins et al. Deep learning for audio signal processing. *IEEE J. Sel. Topics Signal Process.*, 2019. [arXiv:1905.00078](https://arxiv.org/abs/1905.00078).
- [8] D. Stoller, S. Ewert, and S. Dixon. Wave-U-Net: A multi-scale neural network for end-to-end audio source separation. In *ISMIR*, 2018. [arXiv:1806.03185](https://arxiv.org/abs/1806.03185).
- [9] C. Macartney and T. Weyde. Improved speech enhancement with the Wave-U-Net. [arXiv:1811.11307](https://arxiv.org/abs/1811.11307), 2018.
- [10] A. Jansson et al. Singing voice separation with deep U-Net convolutional networks. In *ISMIR*, 2017. <https://openaccess.city.ac.uk/id/eprint/19289/1/7bb8d1600fba70dd79408775cd0c37a4ff62.pdf>.
- [11] Y. Luo and N. Mesgarani. Conv-TasNet: Surpassing ideal time-frequency magnitude masking for speech separation. *IEEE/ACM TASLP*, 2019. [arXiv:1809.07454](https://arxiv.org/abs/1809.07454).
- [12] Y. Luo, Z. Chen, and T. Yoshioka. Dual-path RNN: Efficient long sequence modeling for time-domain speech separation. In *ICASSP*, 2020. [arXiv:1910.06379](https://arxiv.org/abs/1910.06379).
- [13] O. Ronneberger, P. Fischer, and T. Brox. U-Net: Convolutional networks for biomedical image segmentation. In *MICCAI*, 2015. [arXiv:1505.04597](https://arxiv.org/abs/1505.04597).
- [14] K. He, X. Zhang, S. Ren, and J. Sun. Deep residual learning for image recognition. In *CVPR*, 2016. [arXiv:1512.03385](https://arxiv.org/abs/1512.03385).
- [15] A. Vaswani et al. Attention is all you need. In *NeurIPS*, 2017. [arXiv:1706.03762](https://arxiv.org/abs/1706.03762).
- [16] T. Elsken, J. H. Metzen, and F. Hutter. Neural architecture search: A survey. *JMLR*, 2019. [arXiv:1808.05377](https://arxiv.org/abs/1808.05377).
- [17] B. Zoph and Q. V. Le. Neural architecture search with reinforcement learning. In *ICLR*, 2017. [arXiv:1611.01578](https://arxiv.org/abs/1611.01578).
- [18] H. Pham et al. Efficient neural architecture search via parameter sharing. In *ICML*, 2018. [arXiv:1802.03268](https://arxiv.org/abs/1802.03268).
- [19] J. Schulman et al. Proximal policy optimization algorithms. [arXiv:1707.06347](https://arxiv.org/abs/1707.06347), 2017.
- [20] E. Real et al. Large-scale evolution of image classifiers. In *ICML*, 2017. [arXiv:1703.01041](https://arxiv.org/abs/1703.01041).
- [21] E. Real, A. Aggarwal, Y. Huang, and Q. V. Le. Regularized evolution for image classifier architecture search. In *AAAI*, 2019. [arXiv:1802.01548](https://arxiv.org/abs/1802.01548).
- [22] M. Jaderberg et al. Population based training of neural networks. [arXiv:1711.09846](https://arxiv.org/abs/1711.09846), 2017.
- [23] S. Khadka and K. Tumer. Evolution-guided policy gradient in reinforcement learning. In *NeurIPS*, 2018. [arXiv:1805.07917](https://arxiv.org/abs/1805.07917).
- [24] N. Hansen. The CMA evolution strategy: A tutorial. [arXiv:1604.00772](https://arxiv.org/abs/1604.00772), 2016.
- [25] N. Shazeer et al. Outrageously large neural networks: The sparsely-gated mixture-of-experts layer. In *ICLR*, 2017. [arXiv:1701.06538](https://arxiv.org/abs/1701.06538).
- [26] J. Snoek, H. Larochelle, and R. P. Adams. Practical Bayesian optimization of machine learning algorithms. In *NeurIPS*, 2012. [arXiv:1206.2944](https://arxiv.org/abs/1206.2944).
- [27] M. Tan et al. MnasNet: Platform-aware neural architecture search for mobile. In *CVPR*, 2019. [arXiv:1807.11626](https://arxiv.org/abs/1807.11626).
- [28] C. Liu et al. Progressive neural architecture search. In *ECCV*, 2018. [arXiv:1712.00559](https://arxiv.org/abs/1712.00559).
- [29] S. Wisdom et al. Unsupervised sound separation using mixture invariant training. In *NeurIPS*, 2020. [arXiv:2006.12701](https://arxiv.org/abs/2006.12701).
- [30] C. Subakan, M. Ravanelli, S. Cornell, M. Bronzi, and J. Zhong. Attention is all you need in speech separation. In *ICASSP*, 2021. [arXiv:2010.13154](https://arxiv.org/abs/2010.13154).
- [31] Y. Gong, Y.-A. Chung, and J. Glass. AST: Audio spectrogram transformer. In *INTERSPEECH*, 2021. [arXiv:2104.01778](https://arxiv.org/abs/2104.01778).
- [32] C. Finn, P. Abbeel, and S. Levine. Model-agnostic meta-learning for fast adaptation of deep networks. In *ICML*, 2017. [arXiv:1703.03400](https://arxiv.org/abs/1703.03400). (Screened.)

- [33] H. Liu, K. Simonyan, and Y. Yang. DARTS: Differentiable architecture search. In *ICLR*, 2019. [arXiv:1806.09055](#). (*Screened.*)
- [34] A. Défossez et al. Demucs: Deep extractor for music sources. [arXiv:1911.13254](#), 2019. (*Screened.*)
- [35] A. Défossez, G. Synnaeve, and Y. Adi. Real time speech enhancement in the waveform domain. In *INTERSPEECH*, 2020. (*Screened.*) [arXiv:2006.12847](#).
- [36] A. Défossez. Hybrid spectrogram and waveform source separation. [arXiv:2111.03600](#), 2021. (*Screened.*)
- [37] Z. Kong et al. DiffWave: A versatile diffusion model for audio synthesis. In *ICLR*, 2021. [arXiv:2009.00713](#). (*Screened.*)
- [38] S. Pascual, A. Bonafonte, and J. Serrà. SEGAN: Speech enhancement generative adversarial network. In *INTERSPEECH*, 2017. (*Screened.*) [arXiv:1703.09452](#).